

Priority Based Scheduling in Queuing Systems

PMCS Project

Team 4

Problem Statement in Brief




Figure 31.1. When a server frees up, it takes the job at the head of the highest priority, non-empty queue.

- ▶ Queues covered in class followed the FCFS scheduling policy.
- ▶ We shall explore priority-based scheduling where priorities are assigned according to job size (shortest job has the highest priority). We are assuming 10 classes of jobs.
- ▶ Each class is an imaginary M/G/1 queue. When the server becomes free, it always chooses the job at the head of the highest priority non-empty queue to work on.

Problem Statement in Brief

- ▶ How will such a queue be modelled as a CTMC? What is its stationary distribution?
- ▶ What are conditions that ensure queue is stable?
- ▶ What is the distribution for waiting time or sojourn time for the system?
- ▶ What are the average performance metric formulas such as the mean sojourn time, mean number of jobs in the system, or blocking probabilities?
- ▶ Simulations of the queuing system to verify the results derived

Notations

- ▶ S_k = size of priority k job
- ▶ $E[N_Q(k)]$ = average number of priority k jobs in the queue
- ▶ $E[T_Q(k)]$ = average time in queue for priority k jobs
- ▶ $E[T(k)]$ = average time in system for priority k jobs
- ▶ $\lambda_k = \lambda \cdot p_k$ = average arrival rate of jobs of priority k
- ▶ $\rho_k = \lambda_k \cdot E[S_k]$ = contribution to the load due to jobs of priority k
- ▶ $\sum_{i=1}^n \rho_i = \sum_{i=1}^n \lambda_i E[S_i] = \sum_{i=1}^n \lambda_i p_i E[S_i] = \lambda E[S] = \rho$
- ▶ ρ is the server utilization (< 1)

Notations

- ▶ Let S denote size of arbitrary job
- ▶ $E[S] = \sum_{k=1}^n \rho_k E[S_k]$
- ▶ $E[S^2] = \sum_{k=1}^n \rho_k E[S_k^2]$
- ▶ Using Hitchhiker's Paradox, $E[S_e] =$ mean remaining time of a job $= \frac{E[S^2]}{2E[S]}$

Deriving $T_Q(k)$ -Time in Queue for Jobs of Priority k

- ▶ Consider a priority k arrival. That arrival has to wait for
 - ▶ The job currently in service, if there is one.
 - ▶ All jobs of priority $1, 2, \dots, k$ in queue when the job arrives.
 - ▶ All jobs of priority $1, 2, \dots, k-1$ that arrive while the new job is waiting.

- ▶ It can be shown that

$$E[T_Q(k)]^{NP-Priority} = \frac{\rho E[S_e]}{(1 - \sum_{i=1}^k \rho_i)(1 - \sum_{i=1}^{k-1} \rho_i)}$$

Deriving $T_Q(1)$ -Time in Queue for Jobs of Priority 1

- ▶ Consider a priority 1 arrival. That arrival has to wait for
 - ▶ The job currently in service, if there is one.
 - ▶ All jobs of priority 1 in queue when the job arrives.
- ▶
$$\begin{aligned} E[T_Q(1)] &= P(\text{Server busy}) \cdot E[S_e] + E[N_Q(1)] \cdot E[S_1] \\ &= \rho E[S_e] + E[T_Q(1)] \cdot \lambda_1 \cdot E[S_1] \\ &= \rho E[S_e] + E[T_Q(1)] \cdot \rho_1 \\ &= \frac{\rho E[S_e]}{1 - \rho_1} \end{aligned}$$

Deriving $T_Q(2)$ -Time in Queue for Jobs of Priority 2

- ▶ Consider a priority 2 arrival. That arrival has to wait for
 - ▶ The job currently in service, if there is one.
 - ▶ All jobs of priority 1 in queue when the job arrives.
 - ▶ All jobs of priority 1 that arrive while the new job is waiting (not in service).

$$\begin{aligned} \text{▶ } E[T_Q(2)] &= \rho E[S_e] + E[N_Q(1)] \cdot E[S_1] + E[N_Q(2)] \cdot E[S_2] + \\ &\quad E[T_Q(2)] \cdot \lambda_1 E[S_1] \\ &= \rho E[S_e] + E[T_Q(1)] \cdot \rho_1 + E[T_Q(2)] \cdot \rho_2 + E[T_Q(2)] \cdot \rho_1 \\ E[T_Q(2)] \cdot (1 - \rho_1 - \rho_2) &= \rho E[S_e] + \rho_1 \cdot E[T_Q(1)] \\ E[T_Q(2)] &= \frac{\rho E[S_e]}{(1 - \rho_1)(1 - \rho_1 - \rho_2)} \end{aligned}$$

- ▶ Similarly using induction, we can derive for formula
$$E[T_Q(k)]^{NP-Priority} = \frac{\rho E[S_e]}{(1 - \sum_{i=1}^k \rho_i)(1 - \sum_{i=1}^{k-1} \rho_i)}$$

Interpreting the formula

- ▶ $E[T_Q(k)]^{NP-Priority} = \frac{\rho E[S_e]}{(1 - \sum_{i=1}^k \rho_i)(1 - \sum_{i=1}^{k-1} \rho_i)}$
- ▶ $\rho E[S_e]$ is due to waiting for the job in service
- ▶ $(1 - \sum_{i=1}^k \rho_i)$ due to waiting for jobs in the queue of higher or equal priority
- ▶ $(1 - \sum_{i=1}^{k-1} \rho_i)$ is due to those jobs that arrive after our job, but have strictly higher priority than our tagged job
- ▶ This is the average waiting time for a job of class k. How do we obtain the average waiting time across all classes
$$E[T_Q]^{NP-Priority} = \sum_{k=1}^n E[T_Q(k)] \cdot p_k = \sum_{k=1}^n E[T_Q(k)] \cdot \frac{\lambda_k}{\lambda}$$

Shortest-Job-First

- ▶ Consider again the situation of n priority classes. Let's assume that the job sizes range between $x_0 = 0$ and x_n . Job sizes are generated as instances of exponential random variable(μ).
- ▶ Define boundary points x_1, x_2, \dots, x_n such that $x_0 < x_1 < x_2 < \dots < x_{n-1} < x_n$.
- ▶ Assign all jobs of size (x_{k-1}, x_k) to class k .




Figure 31.2. Defining classes based on job size.

Shortest-Job-First

- ▶ Now consider the situation where $n \rightarrow \infty$ and $(x_k - x_{k-1}) \rightarrow 0$. That is, the number of classes, n , is allowed to grow to ∞ , in such a way that $(x_k - x_{k-1})$ becomes arbitrarily small.
- ▶ We are interested in the expected waiting time for a job of size x_k

Expected Waiting Time in SJF

- ▶ Load of jobs in class 1 to k = Load of jobs of size $< x_k = \lambda \int_{t=0}^{x_k} tf(t)dt$
- ▶ Load of jobs in class 1 to $k-1$ = Load of jobs of size $< x_{k-1} = \lambda \int_{t=0}^{x_{k-1}} tf(t)dt = \lambda \int_{t=0}^{x_k} tf(t)dt$
- ▶ $E[T_Q(x)]^{SJF} = \frac{\rho E[S_e]}{(1 - \lambda \int_{t=0}^x tf(t)dt)^2}$
- ▶ $E[T_Q]^{SJF} = \int_{x=0}^{x_n} E[T_Q(x)]f(x)dx = \rho E[S_e] \int_{x=0}^{x_n} \frac{f(x)dx}{(1 - \lambda \int_{t=0}^x tf(t)dt)^2}$

Motivation

- ▶ Scheduling policies can be categorized into non-preemptive and preemptive types.
- ▶ Non-preemptive policies exhibit drawbacks, particularly in scenarios with highly variable job sizes, leading to longer mean response times.
- ▶ Preemptive policies offer better performance in such scenarios by interrupting larger jobs to serve smaller ones, reducing response times.
- ▶ Policies utilizing job size or age can further improve performance by favoring smaller jobs, as seen in the Feedback (FB) scheduling policy.

Preemptive Priority Queueing

- ▶ We continue with the assumptions from NP-Priority:
 - ▶ There are n classes, with Class 1 having the highest priority.
 - ▶ Class k jobs arrive according to a Poisson process with rate $\lambda_k = \lambda \cdot p_k$.
 - ▶ Class k jobs have service requirements with moments $E[S_k]$ and $E[S_k^2]$.
 - ▶ The load of class k is $\rho_k = \lambda_k \cdot E[S_k]$.
- ▶ Preemptive priority queueing differs from non-preemptive queueing in that whenever a job arrives with a higher priority than the job currently in service, the job in service is preempted, and the higher priority job begins service.
- ▶ No work is lost under preemptions.
- ▶ We will compute $E[T(k)]^{P-Priority}$, the mean time in system for a job of priority k in a system with preemptive priority.

Calculating Mean Time in System for Priority k

- ▶ To compute $E[T(k)]^{P-Priority}$, the mean time in system for a job of priority k in a preemptive priority system, we break it down into three components:
 1. $E[S_k]$: Mean service time for a job of priority class k .
 2. Expected time to complete service on all jobs of priority 1 to k already in the system when our arrival walks in.
 3. Expected total service time required for all jobs of priority 1 to $k - 1$ that arrive before our arrival leaves.

Calculating Component (3)

- ▶ Component (3) is the expected total service time required for all jobs of priority 1 to $k - 1$ that arrive before our arrival leaves.
- ▶ This is calculated as:

$$(3) = \sum_{i=1}^{k-1} E[T(k)] \cdot \lambda_i \cdot E[S_i] = E[T(k)] \sum_{i=1}^{k-1} \rho_i$$

Calculating Component (2)

- ▶ We cannot directly sum up the expected number of jobs in each class for classes 1 to k , weighted by the mean job size for that class, as in non-preemptive priority queueing.
- ▶ We make the following arguments to determine (2):
 1. Expected remaining work in the system due to only jobs of priority 1 through k .
 2. Total expected remaining work in preemptive priority system if the system only ever had arrivals of priority 1 through k .
 3. Total expected remaining work in the system if the system only ever had arrivals of class 1 through k and the scheduling order was any work-conserving order.

Calculating Component (2) - Continued

- Using these arguments, we arrive at:

$$(2) = \frac{\lambda \sum_{i=1}^k p_i E[S_i]}{1 - \sum_{i=1}^k \rho_i} \cdot \frac{\sum_{i=1}^k p_i E[S_i^2]}{2 \sum_{i=1}^k p_i E[S_i]} = \frac{\lambda \sum_{i=1}^k \rho_i E[S_i^2]}{2(1 - \sum_{i=1}^k \rho_i)}$$

$$(2) = \frac{\sum_{i=1}^k \rho_i \frac{E[S_i^2]}{2E[S_i]}}{1 - \sum_{i=1}^k \rho_i}$$

Combining Components

- Finally, adding (1), (2), and (3), we have:

$$E[T(k)]^{P-Priority} = E[S_k] + \frac{\sum_{i=1}^k \rho_i \frac{E[S_i^2]}{2E[S_i]}}{1 - \sum_{i=1}^k \rho_i} + E[T(k)] \sum_{i=1}^{k-1} \rho_i$$

$$E[T(k)](1 - \sum_{i=1}^k \rho_i) = E[S_k] + \frac{\lambda \sum_{i=1}^k \rho_i \frac{E[S_i^2]}{2E[S_i]}}{1 - \sum_{i=1}^k \rho_i}$$

$$E[T(k)]^{P-Priority} = \frac{E[S_k]}{(1 - \sum_{i=1}^k \rho_i)} + \frac{\sum_{i=1}^k \rho_i \frac{E[S_i^2]}{2E[S_i]}}{(1 - \sum_{i=1}^{k-1} \rho_i)(1 - \sum_{i=1}^k \rho_i)}$$

Interpretation of $E[T(k)]^{P-Priority}$

- ▶ For preemptive service disciplines, the time in the system of a job can be divided into two components:
 1. Time until the job first starts serving (waiting time), denoted by *Wait*.
 2. Time from when the job first receives some service until it leaves the system (residence time), denoted by *Res*.

Understanding Mean Time in System for Priority k

- ▶ **Question:** Is residence time the same as service time?
- ▶ **Answer:** No. The residence time is longer as it includes interruptions.
- ▶ **Question:** What does the first term $E[S_k]/(1 - \sum_{i=1}^{k-1} \rho_i)$ in expression $E[T(k)]^{P-Priority}$ represent?
- ▶ **Answer:** This represents the mean residence time of the job of class k , denoted as $E[Res(k)]$. It signifies the expected length of a busy period started by a job of size $E[S_k]$, where only jobs of class 1 through $k - 1$ are allowed in the busy period.

Understanding Mean Time in System for Priority k (Continued)

- ▶ **Question:** What does the second term in the expression $E[T(k)]^{P-Priority}$ means ?.
- ▶ **Answer:** The remaining term in expression $E[T(k)]^{P-Priority}$ represents the mean time until the job of priority k first receives service. It is calculated as:

$$E[Wait(k)] = \frac{\sum_{i=1}^k \rho_i \left(\frac{E[S_{2i}]}{2E[S_i]} \right)}{\left(1 - \sum_{i=1}^{k-1} \rho_i \right) \left(1 - \sum_{i=1}^k \rho_i \right)}$$

This term is similar to $E[T_Q(k)]$ for the non-preemptive priority queue, but only considers interruptions by jobs of class 1 through k . This reflects the fact that a job of class greater than k will be preempted immediately.

Understanding Mean Time in System for Priority k (Continued)

- It is sometimes convenient to rewrite the expression $E[T(k)]^{P-Priority}$ as:

$$E[T(k)]^{P-Priority} = \frac{E[S_k]}{1 - \sum_{i=1}^{k-1} \rho_i} + \frac{\frac{\lambda}{2} \sum_{i=1}^k \rho_i E[S_{2i}]}{(1 - \sum_{i=1}^{k-1} \rho_i)(1 - \sum_{i=1}^k \rho_i)}$$

Comparison with Non-Preemptive Priority Queueing

- ▶ **Question:** Does a high priority job obtain better performance in preemptive priority queueing?
- ▶ **Answer:** Yes. Both terms in $E[T(k)]^{P-Priority}$ depend only on the first k priority classes. Thus, a high priority job in preemptive priority queueing indeed performs better, even in a high-variability job size distribution, compared to non-preemptive priority queueing.
- ▶ **Note:** In the non-preemptive case, the expression for $E[T_Q(k)]^{NP-Priority}$ has a similar denominator, but the numerator involves all n classes contributing to the excess, unlike in the preemptive case.

Preemptive Shortest Job First(PSJF)

- ▶ A preemption only occurs when a new job arrives whose size is smaller than the original size of the job in service.

- ▶ Recall that

$$E[T(k)]^{P-Priority} = \frac{E[S_k]}{1 - \sum_{i=1}^{k-1} \rho_i} + \frac{\frac{\lambda}{2} \sum_{i=1}^k \rho_i E[S_i^2]}{(1 - \sum_{i=1}^{k-1} \rho_i)(1 - \sum_{i=1}^k \rho_i)}$$

- ▶ Performing the same limiting operations as we did in analyzing SJF (where we imagine that jobs of size x form one “class” and that there are an infinite number of classes), we get

$$E[T(x)]^{PSJF} = \frac{x}{1 - \rho_x} + \frac{\frac{\lambda}{2} \int_0^x f(t) t^2 dt}{(1 - \rho_x)^2}$$

- ▶ Here $f(t)$ is the p.d.f. of job size, S and $\rho_x = \lambda \int_0^x t f(t) dt$ is defined to be the load made up by jobs of size less than x

Preemptive Shortest Job First(PSJF)- Busy Period Analysis

- ▶ We start by breaking up response time into waiting time and residence time
- ▶ $E[T(x)]^{PSJF} = E[Wait(x)]^{PSJF} + E[Res(x)]^{PSJF}$
- ▶ $Res(x)$ is just the duration of a busy period started by a job of size x , where the only jobs that make up this busy period are jobs of size $\leq x$.
$$E[Res(x)]^{PSJF} = \frac{x}{1-\rho_x}$$

Preemptive Shortest Job First(PSJF)- Busy Period Analysis



Figure 32.1. Transformer glasses for PSJF. Whereas in FB, the transformer glasses truncate all jobs of size $> x$ to size x , in PSJF, the transformer glasses make jobs of size $> x$ invisible.

- ▶ Let us now calculate $E[Wait(x)]^{PSJF}$
- ▶ For a job of size x , busy period is the job of sizes $\leq x$. Let us call this W_x
- ▶ We will use S_x to denote the size of a job of size $\leq x$. The density of S_x is $\frac{f(t)}{F(x)}$ where $f(t)$ is density of S .


Preemptive Shortest Job First(PSJF)- Busy Period Analysis

$$\begin{aligned}\blacktriangleright E[Wait(x)]^{PSJF} &= \frac{E[W_x]}{1-\rho_x} \\ &= \frac{E[T_Q|S_x]^{FCFS}}{1-\rho_x} \\ &= \frac{\frac{\lambda F(x)E[S_x^2]}{2(1-\rho_x)}}{1-\rho_x} \\ &= \frac{\lambda F(x) \int_0^x t^2 \frac{f(t)}{F(x)} dt}{2(1-\rho_x)^2} \\ &= \frac{\lambda \int_0^x t^2 f(t) dt}{2(1-\rho_x)^2}\end{aligned}$$

$$\blacktriangleright \text{Thus, } E[T(x)]^{PSJF} = \frac{x}{1-\rho_x} + \frac{\lambda \int_0^x t^2 f(t) dt}{2(1-\rho_x)^2}$$

Simulations

Simulations- Details of the System



- ▶ We have implemented a M/M/1 queue where arrival rate $= \text{Exp}(\lambda) = 0.1$ and service requirements, $S_x = \text{Exp}(\mu) = 0.2$
- ▶ The system has 4 queues and 1 server. Each queue is for a single priority. Priority is assigned according to job sizes.

Preemptive SJF-Theoretical Results

$$\blacktriangleright E[T_Q]^{SJF} = \int_{x=0}^{x_n} E[T_Q(x)]f(x)dx = \\ E[S_e] \int_{x=0}^{x_n} \frac{f(x)dx}{(1-\lambda \int_{t=0}^x tf(t)dt)^2}$$

$$\blacktriangleright E[T_Q]^{SJF} = \int_{x=0}^{x_n} E[T_Q(x)]f(x)dx = \\ \frac{\rho E[S^2]}{2E[S]} \int_{x=0}^{x_n} \frac{f(x)dx}{(1-\lambda \int_{t=0}^x tf(t)dt)^2}$$

$$\blacktriangleright \text{Substituting values, } E[T_Q]^{SJF} = \frac{\rho E[S^2]}{2E[S]} \int_{x=0}^{x_n} \frac{f(x)dx}{(1-\lambda \int_{t=0}^x tf(t)dt)^2}$$

Non-Preemptive: Theoretical Calculations

- ▶ Given parameters:

$$\lambda = 0.1$$

$$E[S] = \frac{1}{0.2} = 5$$

$$\rho = \lambda \cdot E[S] = 0.5$$

$$E[S^2] = \frac{1}{(0.2)^2} + \frac{1}{0.04} = 50$$

$$\frac{\rho \cdot E[S^2]}{2 \cdot E[S]} = 2.5$$

Theoretical Calculations (Continued)


- ▶ We calculate the integral:

$$\int_0^{\infty} \frac{e^{-0.2x}}{(1 - (0.5 - (0.5x + 2.5)e^{-0.2x}))^2} dx = 1.45231388191$$

- ▶ Then, we multiply it by $\frac{\rho \cdot E[S^2]}{2 \cdot E[S]}$:

$$\frac{\rho \cdot E[S^2]}{2 \cdot E[S]} \cdot \int_0^{\infty} \frac{e^{-0.2x}}{(1 - (0.5 - (0.5x + 2.5)e^{-0.2x}))^2} dx = 3.63078470$$

Results obtained through Simulations



Avg waiting time = 3.669

- Average Waiting Time: 3.669

Conclusions

- ▶ Both theoretical and simulation results match, indicating the accuracy of the theoretical model in predicting the average waiting time.

Preemptive: Theoretical Calculations

- ▶ Given parameters:

$$\lambda = 0.1$$

$$\mu = 0.2$$

Theoretical Calculations (Continued)

- First let's calculate loads made up by jobs each queue less than x :

$$\rho_x = \lambda \cdot \int_0^x tf(t)dt$$

- We have four bins $[0, 2.5]$, $[2.5, 5]$, $[5, 7.5]$, $[7.5, \infty)$:

$$\rho_1 = \lambda \cdot \int_0^{2.5} tf(t)dt$$

$$\rho_1 = \lambda \cdot \int_0^{2.5} x \cdot \mu e^{-\mu \cdot x} dx$$

$$\rho_1 = 0.04$$

$$\rho_2 = \lambda \cdot \int_0^5 tf(t)dt$$

$$\rho_2 = \lambda \cdot \int_0^5 x \cdot \mu e^{-\mu \cdot x} dx$$

$$\rho_2 = 0.13$$

Theoretical Calculations (Continued)

$$\rho_3 = \lambda \cdot \int_0^{7.5} tf(t) dt$$

$$\rho_3 = \lambda \cdot \int_0^{7.5} x \cdot \mu e^{-\mu \cdot x} dx$$

$$\rho_3 = 0.22$$

$$\rho_4 = \lambda \cdot \int_0^{\infty} tf(t) dt$$

$$\rho_4 = \lambda \cdot \int_0^{\infty} x \cdot \mu e^{-\mu \cdot x} dx$$

$$\rho_4 = 0.5$$

Theoretical Calculations (Continued)



$$E[T(x)]^{\text{PSJF}} = \frac{x}{1 - \rho_x} + \frac{\lambda \int_0^x t^2 f(t) dt}{2(1 - \rho_x)^2}$$

Using the above formula we get the following values:

$$E[T(2.5)] = 1.08$$

$$E[T(5)] = 1.414$$

$$E[T(7.5)] = 2.065$$


$$E[T(\infty)] = 12$$

Preemptive SJF-Theoretical Results

- ▶ Average waiting time =

$$\begin{aligned} & E[T(2.5)] \cdot F_x(2.5) + E[T(5)] \cdot [F_x(5) - F_x(2.5)] \\ & + E[T(7.5)] \cdot [F_x(7.5) - F_x(5)] + E[T(\infty)] \cdot (1 - F_x(7.5)) \\ & = 0.432 + 0.4242 + 0.1652 + 2.64 \\ & = 3.6614 \end{aligned}$$

Results obtained through Simulations



Average waiting time = 3.7

- Average Waiting Time: 3.7

Conclusions

- ▶ Both theoretical and simulation results match, indicating the accuracy of the theoretical model in predicting the average waiting time.